Chapter 8: Conclusion

Luke Elson

Abstract. I briefly recapitulate the main claims of the book.

I've argued that by taking seriously the possibility that our beliefs and desires are shot through with vagueness—that is, indeterminacy or uknowability—in a tolerance-denying way then we can explain a number of otherwise-troubling practical problems and puzzles. I've defended a view of rationality as fundamentally maximising, where this is extended to consider determinacy in whether we maximise expected utility.

There are three main distinctive ideas in the book.

First, Supersharp. Our beliefs and preferences can be vague and thus unsharp, even on a revealed-preference account of them. If the correct account of vagueness is a tolerance-denying one—where tolerance principles such as 'adding 1mm to a short man always leaves you with a short man' are false—then we can reject tolerance principles for preferences and beliefs too:

Supersharp. There are agents whose utility and credence functions are vague (and thus soft-unsharp), nevertheless it is determinately true that they have full, point-valued credences and preferences.

I use *unsharpness* as an umbrella term for incomplete preferences and imprecise credences, and central claim of the book is that unsharpness is vagueness. I didn't argue for a tolerance-denying theory of vagueness, but explored the implications of adopting one.

Supersharp says that 'for each proposition there's a number x that is your credence in that proposition' is determinately, knowably true. But it may be indeterminate or unknowable (depending on your theory of vagueness) *which* number it is, in which case your credence in that proposition is imprecise. Similarly, it can be indeterminate or unknowable whether you prefer A to B or B to A, in which case—depending on whether you are sensitive to sweetening of either—you have incomplete preferences between A and B or are simply indifferent between them.

Supersharp allows us to vindicate claims that we (perhaps insofar as we are rational) have full, transitive preferences as supertruths. It strikes a nice middle ground. We don't need to radically alter our decision theory to cope with preference gaps or cycles. But we also don't need to accept that we have counterintuitively sharp preferences and credences.

The second main idea is the decision rule Compound, defended in Chapter 4:

Compound with Necessary Detachment. Do not complete impermissible compound actions. If (i) for some ψ_2, $\phi_1 \& \psi_2$ is permissible, but (ii) $\phi_1 \& \phi_2$ is impermissible, and (iii) you have done ϕ_1, then ϕ_2 is impermissible.

Compound forbids us from completing determinately inadmissible *sequences* of action or 'compound actions'. I argued that it allows us to cope with unsharpness *without* abandoning standard decision theory. We instead extend it to cope with permissibility-facts that arise at the level of more than one action, as can happen under vagueness (and indeed happens in the sorites).

Compound is the heart of my response to criticisms that unsharpness allows agents to be value-pumped, and that we must thus either demand sharpness or go unorthodox. In Adam Elga's famous case where rejecting two sequential bets A and B seems clearly irrational, Compound forbids her from rejecting Bet B if and only if she already rejected Bet A. I argued that under vagueness this doesn't represent an unacceptablly unorthodox departure from consequentialism, because rejecting Bet A alters the status (and consequences) of rejecting Bet B.

The third main idea is the repeating practical sorites. In the clearest case of a practical sorites we trade off a precise goal where each step 'makes a difference' against a vague goal, where each step seems to 'make no difference'. Rationality thus seems to demand we favour the precise goal at every stage. Suppose you wish to stay hairy, because you have a beautiful head of hair. It's so beautiful that you can sell each strand for £1. Rationality seems to compel you to take every pound: the money counts but one hair seems to make no difference to your vague project of staying hairy.

I responded to arguments that such cases engender rationally cyclic preferences. I argued that the Puzzle of the Self-Torturer and similar cases are practical sorites, and there are no genuinely rational intransitive preferences at work. I distinguished between *single-threshold* and *repeating* practical sorites, arguing that the Self-Torturer faces the latter. (Your hairy self faces the former, and is in a better position.)

The final chapters extended the view in two directions. Chapter 6 expanded repeating practical sorites to *groups*. I focused on the case of climate change which Andreou has correctly argued is analogous to the Puzzle of the Self-Torturer. I argued pessimistically that vagueness will make it difficult for governments—especially democracies—to enforce meaningful restictions on emissions. Chapter 7 turned to spectrum arguments that impartial betterness is intransitive. I argued that they too can be modelled as a sorites, especially as a repeating sorites.

Sometimes I've worried that in defending this package of views where so much turns out to be vagueness, I've been a man with a hammer seeing everything as a nail. That is a risk of course, but the phenomena I've considered do have striking structural similarities, and I've argued that this is because all involve vagueness. Even if the models I've presented fail for one reason or another, it was still worth it to see how far a broadly orthodox view can be pushed when supplemented with vagueness.