

SPECIAL ISSUE ARTICLE OPEN ACCESS

Moral Risk in Marketised Medicine

Luke Elson 

Philosophy Department, University of Reading, Reading, UK

Correspondence: Luke Elson (luke.elson@reading.ac.uk)

Received: 17 June 2025 | Revised: 18 June 2025 | Accepted: 13 July 2025

ABSTRACT

We hope that doctors will recommend and provide the most appropriate investigations and treatments. I argue that some ways of structuring medical provision—mostly, those involving markets—impose a risk of *overprovision*. This is bad financially, medically, and epistemically, and therefore morally. We should be extremely cautious about anything that might damage trust in doctors. Common “who cares if the doctor works for a private company so long as treatment is free?” defences of healthcare marketisation and privatisation miss this important point.

Here is the *British Medical Journal*:

Case 1 questions whether Dr Brown Bear is an unscrupulous private practitioner, rather than an NHS GP. In this case of a probable viral rash, he could reasonably have encouraged self management (with appropriate safety netting) or asked the family to attend surgery for assessment. His decision not to suggests a potential financial incentive for conducting an arguably clinically inappropriate home visit.¹

The BMJ's annual joke issue here takes a look at the ecosystem of the excellent children's cartoon *Peppa Pig*. The passage contains a mixture of implicit and explicit claims: in a viral rash, Dr. Brown Bear's response—a home visit—was clinically inappropriate; such an inappropriate response suggests a financial incentive to avoid a more appropriate course; acting on such a financial incentive is unscrupulous; anyone who acts in that way is (more) likely to be a private practitioner, rather than working through the National Health Service (NHS).

These claims are all about money, with the possible exception of the first. If a home visit is clinically inappropriate because it is

an excessive use of medical resources for something so minor, then that is about money too.

With minor exceptions, NHS treatment has been free at the point of use in the UK since 1948. The four nations—England, Wales, Scotland and Northern Ireland—have had health services for decades, but my focus is England. The NHS is one of the most socialised healthcare systems in the developed world, but since the 1980s, England has had a ‘purchaser-provider split’ between those who hold the money (purchasers or commissioners) and those who provide the care.

The BMJ's joke would make no sense if its background assumptions were not shared or at least understood by the intended audience. Clearly one such assumption is that NHS General Practitioners (primary care doctors or GPs) are less likely to be ‘unscrupulous’ than private (non-NHS) practitioners. But why might this be so? Even if the *British Medical Journal* has fallen into typical ‘envy of the world’ nationalism, what is the basis of that nationalist belief?

The obvious answer is the NHS primary care contract. GPs are not employed by the state as many hospital doctors are, but work for independent ‘practices’ or ‘surgeries’. Some GPs are equity partners with a stake in their practice, and others are employees

This is an open access article under the terms of the [Creative Commons Attribution](https://creativecommons.org/licenses/by/4.0/) License, which permits use, distribution and reproduction in any medium, provided the original work is properly cited.

© 2025 The Author(s). *Ratio* published by John Wiley & Sons Ltd.

of those partners. Typically GP surgeries have been private businesses with only one customer (the NHS) for core medical services. That customer has enormous power and offers a standard GP contract. Take it or leave it.²

Under that contract, care is free at the point of access and payment is required only for prescription drugs, at a current (January 2025) flat rate of £9.90 per item or £114.50 per year (the fee does not go to the GP). English patients might have to open their wallets at the GP surgery for things outside of the NHS contract such as medical fitness letters or travel vaccinations, but not for core medical services unless they wish to 'go private' for faster treatment or longer appointments, for example.

My focus is on the financial incentives for medical care. An NHS GP will either be paid a 'capitation' rate per patient on the books, or a salary or hourly rate if they are an employee. She lacks a financial incentive to visit in the middle of the night and is largely insulated from market forces, so the opportunity to be unscrupulous through *excessive* home visits is unlikely to arise. Conversely, a private GP working outside the NHS may have a price list including—quite reasonably—a high fee for a night-time home visit.

In an example familiar to many of us and to which I will return, the capitation payment is akin to anonymous marking of undergraduate work: even if you would like to unethically down-mark the students you dislike, with anonymous marking you will not usually get the chance.

I will argue that it is morally good when doctors are insulated from market forces. Because of the special features of medical care, it is insidiously morally bad when medical decisions are made by those with a financial stake in those decisions—even if those decisions are never *in fact* influenced by those stakes. Thus this paper considers the 'organisational ethics' of the NHS, and I have been heavily influenced by the work of Lucy Frith, who also considers the possibility that 'something is lost if medicine becomes subject to the same kinds of commercial ethos and pressures as more market-orientated services' (Frith 2018, 52).

In this paper, I will consider a particular way in which something could be lost. Much of my argument hangs on empirical claims about how patients react to various situations (such as liking 'thorough' doctors). I think they are very plausible, but my argument is conditional on (i.e., hostage to) those claims.

1 | Doctors Have Incentives

Doctors get paid, in money, housing, training, whatever. Even holding fixed the amount of money in the system, and user fees, and so on, there are different ways that money can flow *within* the healthcare system. But ultimately money flows from the patient, taxpayer or insurance company to doctors (and to pharmacists, nurses, cleaners, drug companies, hospital construction companies...).

My focus here is on doctors, because I am concerned with money's impact on medical diagnosis and decision-making. But what I say can also apply to nurse practitioners, triage nurses, dentists,

opticians, midwives, and so on, insofar as they do those things. Here are some models for paying doctors:

1. Salaried: Doctors are paid a flat-rate salary for a specified role, not dependent on performance (unless they do so badly as to risk their jobs). Pay is 'all or nothing', with options for working part-time and so on. Imagine a salaried employee at a large NHS hospital.
2. Hourly wage: Doctors are paid by the hour, and it is (partly) up to them to determine their working hours. Locum or bank doctors in the UK work on this basis: they sign up to work shifts at a given rate.
3. Fee for service: Doctors have a menu of services, with prices. This is the model of much non-NHS medical care in the UK, and many hospital doctors in the USA (and partly explains the extraordinary complexity of American medical bills).
4. Capitation: Doctors or groups of them get a fixed annual 'capitation' payment for each registered patient. Whatever is not spent on patient care, they keep. This is the model of General Practice as I sketched it above.

Many doctors mix and match these models. NHS doctors often do 'private work' on the side, offering fee-for-service private care to top up their salary. The idea behind BMJ's joke—which I am now thoroughly stripping of humour by overexplaining it—seems to be that a salaried or capitated NHS doctor has no financial incentive to visit someone's house in the middle of the night, whereas a doctor paid an hourly wage or fee for service does, especially if she can charge more at night or for home care.

It might be claimed that there is no per se ethical difference between these payment models. Especially if there is universal healthcare provision, the choice of payment structure could be seen as fundamentally economic: which is the most efficient?

And there are good moral reasons to choose the most economically-efficient healthcare system. Doing so allows for more complete provision at a lower cost. If the NHS can offer existing services more efficiently, the savings can be spent on treating further ailments. Insofar as relieving suffering is a good, we have reason to do this, and reasons of fairness too, since expanding NHS provision will particularly benefit those who can not afford to 'go private' for treatment that is not available on the NHS, or is available only with such a long waiting list that it might as well not be. Governments have plenty of other things to spend money on too. As well as other treatments, savings in healthcare might allow for better roads or tax cuts.

Thus we can accept that there are strong overarching moral reasons to spend money as efficiently as possible. But my opponent goes further, and claims that this is the only way in which the choice of healthcare payments structures is morally-laden. Beyond considerations of efficiency, the claim goes, it does not matter whether the NHS directly employs all its own doctors or shrinks to a rump of 'commissioning' bureaucrats who negotiate and sign contracts with private hospitals.

That is my opposition. Now of course savings *could* be entirely benign, for example if NHS contracts incentivise doctors to

offer preventative care that is cheaper in the long run. In the news in early 2025 is a government plan to incentivise GPs with £20 to contact specialist hospital doctors for ‘advice and guidance’ rather than referring patients for consultation with said specialists.³

Marketisation could also in-principle bring savings from administrative efficiencies. But they are likely to be marginal if they even exist, considering the cost of contract tendering and negotiations, for example. All such costs are less impressive when set against those of training and paying for doctors and all the assistance and equipment they need.

I will try to set aside economic efficiency and focus on the incentives doctors face. I will now describe in the most cynical terms the incentives engendered by the four payment models described above.

Salaries incentivise doctors to do the minimum amount of work needed to keep their jobs (or get promoted, or similar), and to choose a post and patients that pay the highest salary for the least effort. Depending on its level, an hourly wage incentivises doctors either to work just enough to pay the bills, or—at the other end of the scale—to work past the point of dangerous fatigue and maximise income. (In the worst case, the doctor must work past the point of dangerous fatigue to pay the bills.)

Fees for service offer direct incentives to provide some investigations or treatments rather than others. Different items on the menu bring in different amounts of money.

Capitation does not allow the doctor to directly influence how much money she receives for a given patient population. The incentive is to choose the cheapest treatment that will keep the patient alive and on the books or, in extreme cases—if a certain patient can *never* be profitable—off the books in one way or another. There is also an incentive to recruit a cheap patient pool, for example, through advertising near local universities.

I felt a little treacherous writing those paragraphs, especially since I have benefitted enormously from many different kinds of medical care since I was born (and before). But doctors are not saints, and we routinely worry about perverse incentives tempting financial advisors and car mechanics, so why would medicine be any different? It seems horribly naive to claim that the choice of medical payment system—and thus of the incentives doctors face—is *merely* a matter of efficiency. The effectiveness of financial incentives is also recognised in the new English payments to promote Advice and Guidance, for example. Why offer the money if it will not affect medical behaviour?

2 | Three Dangers of Medical Care

But what exactly are we paying doctors for? The central tasks of a medical consultation are to assess a patient’s clinical need and recommend a course of action. That course could be further investigations or tests, treatments, referrals to another doctor, ‘wait and see’, or perhaps nothing.

I will assume that each (potential) patient has a clinical need, which may be nil if they are perfectly healthy. Clinical need is for me an *ideal advisor* notion: yours is what the best doctor acting purely in your interests with no financial or other constraints would recommend for you. Such a notion is clearly an idealisation—there might be reasonable disagreement between advisors, for example—but it is enough to be getting on with.

We hope for *correct provision*, when the medical care proposed or provided matches the clinical need. If we are to agree with the BMJ, ‘self management (with appropriate safety netting)’ would have been correct provision in Case 1. In correct provision, the doctor correctly identifies and executes the best course of action for the patient, focusing purely on what is medically best for that patient.

But that is morally uninteresting compared to the alternatives. There would seem to be three other possibilities: the patient receives insufficient care, excessive care, or the wrong care.

In *underprovision*, the patient is denied needed care. Moral and political worries about healthcare understandably focus on underprovision, especially where people cannot afford the care they need. But we can distinguish two kinds of underprovision you might face. Sometimes there simply is not enough money in the system:

Supply underprovision. There are inadequate resources entering the healthcare system, and for this reason a patient’s clinical need isn’t met.

Supply underprovision is obviously not ideal, but it need not imply an ethical problem or that anyone is acting wrongly. As discussed above, resources are finite and we must spend money on things other than healthcare. Hence the moral reason I mentioned above to choose the most economically-efficient healthcare system: if the relief of suffering is a moral good, then we should get as much of it as we can with the money we have.

The NHS clearly faces supply underprovision. This need not imply unfairness *within* the healthcare system; such underprovision can happen even as the healthcare system is doing the best, fairest job it can with the resources it has.

Supply underprovision is where questions of rationing emerge: if the healthcare system can not meet everyone’s clinical need, some must receive underprovision. And there’s particular controversy over *bedside rationing* where such decisions are made by the doctor, with identifiable patients. On the face of it, bedside rationing conflicts with the doctor’s duty to seek the best care for each of her patients, considered as individuals.

Another kind of underprovision *does* seem unfair, and even in the context of rationing we should seek to avoid it:

Distributive underprovision. A patient’s clinical need is not met, even though there are sufficient resources in the healthcare system to meet that need without unfairness.

Imagine that you are denied a lifesaving cancer drug because the healthcare system has spent its money on expensive but ineffective treatments for less-serious diseases. Simply put, the money was there but it went to the wrong place.

I will focus on distributive underprovision amongst patients who in some sense have equal right to medical care, such as NHS patients. The first guiding principle of the NHS Constitution for England is ‘a comprehensive service, available to all’, but the second is that ‘access to NHS services is based on clinical need, not an individual’s ability to pay’.⁴

The NHS explicitly aims to match spending to clinical need via the number of quality-adjusted life years (QALY) that a treatment can be expected to add. There is no bright-line threshold, but in general, a treatment that costs under £20,000 per QALY is likely to be considered cost-effective and thus funded, whereas one that costs more than £30,000 per QALY is unlikely to be.⁵ Insofar as clinical need tracks QALYs, these criteria aim to avoid distributive underprovision. I think the moral badness of distributive underprovision is clear, so I will not focus on it.

Even outside universal healthcare systems, depending on the terms and conditions, all of those holding a particular insurance policy at the same tier have a right to be treated equally, according to their clinical needs. Nevertheless, my focus is on the NHS, partly for simplicity (it lets me avoid the problem of there being many insurance companies with differing policies) and partly because I want to argue that the problem I am exploring can arise under universal healthcare too.

Holding fixed the efficiency of the healthcare system, if you have been a victim of distributive underprovision this is because some other patient has received more than their fair share of the medical resources. Where else could the money have gone?

Just as we distinguished between supply and distributive underprovision, so we must between two kinds of overprovision, where the patient receives ‘too much’ care. In supply overprovision this is care in excess of her clinical need. On the other hand, in *distributive overprovision* she receives more than her fair share of healthcare resources.

One can face both supply underprovision and distributive overprovision at the same time, if one gets an unfairly large piece of the pie... but that piece is still too small to meet one’s dietary needs.

Whereas the rationing debate concerns underprovision, I will now argue that the real problem for marketisation is the incentive to overprovide.

3 | Overprovision Is Harmful and Likely

The BMJ suspects Dr. Brown Bear of overprovision, in the form of a clinically unjustified night-time visit. If he is indeed a private provider, things are a little murky here. He can notrealistically be accused of distributive overprovision outside the NHS, but he is engaged in something like supply overprovision either because the night-time visit exceeds the patient’s clinical need—and is

inadvisable even when money is no object—or because any clinical benefit is not worth the financial cost to the patient.

We must distinguish between these cases. Those where healthcare is completely useless and unnecessary and those where it is simply not worth the cost seem quite different. Dr. Brown Bear seems to fall on the latter side of this divide, since he is visiting to examine a rash, which *could* be something serious.

Surely his visit is not outright harmful? And in general, if healthcare is not completely otiose but not worth the cost, how could absolute overprovision be harmful if money is no object? I do not think travelling first-class is worth it, but were money no object then I would buy a first-class ticket. Why not, after all?

In fact, medical overprovision is harmful in many ways. First of all, in this case—assuming the BMJ is right about who he works for—there is the bill. If the clinical benefit was not worth that price, then the patient is worse off in the round, because she is overpaid. Now, it might be argued, just like travelling first-class, that it is the patient’s choice to make.

In many commercial exchanges, such ‘eyes open’ danger to the wallet would be the main problem. And indeed if Dr. Brown Bear warned the patient that he could come out and take a look tonight but it will not be cheap and it can almost certainly wait until the morning, then there is little to object to on financial grounds. We might think both that few can afford to spend the money on the extra ‘luxury’ care, and that those who can afford it know what they are getting into. So just like first-class travel, no harm done.

Bernard Williams argued that that purchased overtreatment (where having enough money is a sufficient condition of medical treatment, and the spendy patient wants treatment *beyond* her clinical need) is likely to be a marginal phenomenon, ‘since only a few hypochondriacs desire treatment when they do not need it’.⁶ This is correct in a very narrow sense: given knowledge that they do not need some treatment, only a few people would desire it.

But I think Williams is too sanguine because of his focus on *treatment* rather than *investigations*, because he overlooks the information-asymmetry between patient and doctor, and because he focuses on cases where the healthcare is outright unnecessary (‘they do not need it’).

Yes indeed, only a few patients desire treatments they know they do not need. But we patients rarely know what we need—and especially rarely know this directly based on the medical evidence, even if we do have Google’s assistance—so we defer to our doctors. Doctors not only provide (some) investigations or treatments, they recommend them, relying on evidence and training that very few patients can match. We do not need to reach for hypochondriacs to consider ways in which overtreatment might arise, or even to draw the analogy with a first-class flyer. We are not like someone with full information weighing the costs and benefits of a first-class ticket.

The situation is more akin to seeing a roofer or a car mechanic. The treatment is expensive, the professional *investigation* is not

cheap (especially if it requires scaffolding or taking your engine apart), and your source of information is that very professional, in a way that you can not easily verify. To even see whether there is a serious problem is non-trivial. You might be able to climb the scaffolding to take a look if you do not mind violating an insurance policy, but what is the point? You do not know what you are looking at. Like putting up scaffolding, medical investigations themselves can be extremely expensive and can be venues for overprovision without hypochondria.

Even without outright deceit, overprovision presents an epistemic danger to the patient. When a doctor recommends a given investigation, she is at the very least implying that, in her clinical judgement, the evidence suggests it is worth the costs (financial and otherwise). In particular, that you may have a disease the test would detect. In the most egregious cases, there is *no* evidence that you have such a disease, but in less extreme versions, a serious disease may be very weakly suggested by the evidence. Being led into false beliefs about what the evidence supports, especially about such intimate matters, is a harm.

In a standard medical cliché, given the evidence or symptoms—an equine silhouette, the clomp of hooves—your disease might be a zebra. But at least in my living situation (in suburban England near to some farms but no zoo) the disease is far more likely to be something less serious and more commonplace. A horse, for example. In cases of underdetermination by data where the evidence is consistent with both a commonplace and a rare phenomenon, the thought is that we should in general attach more credence to the commonplace phenomenon. Nevertheless, Equine Resonance Imaging would settle the matter—by testing for the rare phenomenon—for a price. That price may fall on my wallet, or it may fall on the collective wallet.

And it need not just be financial, though modern medical imaging is very expensive. One must also face possible radiation dangers, time off work, the stress and fear of sitting in an uncomfortable waiting room with bad coffee, and so on. A blood draw or MRI is no fun even in a hotel-like private hospital, and neither is the wait for the results. You might think that none of this matters in a spurious home visit for a rash, but would not it be an anxious wait for the doctor to arrive? It must be serious if he is willing to come out in the middle of the night... he says it is nothing, but maybe he is just trying to put me at ease? This is quite different to the choice of travel class.

A long-term epistemic risk of overinvestigation is that of the boy who cried wolf. When there eventually was a wolf, nobody believed him. A doctor who consistently recommends expensive, inconvenient, or simply painful testing on the basis of (unbeknownst to you) flimsy evidence may similarly not be believed when he recommends you an investigation this time when the evidence is (again, unbeknownst to you) rather stronger. Trust is damaged. Moreover, trust may be damaged in *other* doctors.

I have argued that overinvestigation presents distinctive harms. Investigations are also likely to flag up benign abnormalities, leading to further investigation and perhaps overtreatment.

Investigations may often have some benefits even if they are not worth the cost. If my sore head will probably feel better within

a few days, then a CT scan would be supply overprovision: not useless but far in excess of my clinical need, and I would rather not spend the money or radiation exposure on it. Alternatively, we could have a case of distributive overprovision. Even if there are enough red flags in my case that a CT *would* be sensible for me, there might nevertheless be a backlog of patients who would benefit from it rather more. In that case, for me to receive the scan ahead of them would be unfair.

Williams is correct that very few of us desire *treatment* they outright do not need, but actual decisions are much murkier, especially for investigations. We could be sucked into overprovision, without being one of Williams's 'hypochondriacs'. We might not know we are not ill, and we might wish for effective but not *cost-effective* treatments if we are mildly ill.

I have argued that overprovision, especially of investigations, is harmful. I will now argue that unfortunately, marketisation makes overprovision particularly likely. In short, it is hard to detect and patients like it.

In a market, overprovision is more likely to be profitable than underprovision. Fundamentally, underprovision involves turning away paying customers, given some kind of fee-for-service model. Dr. Brown Bear is unscrupulous because he profits from an overprovided ('clinically inappropriate') home visit. There is no comparable financial incentive for him to underprovide, for example by refusing a home visit when one *would* have been appropriate. Healthcare markets and their incentives are complex, so there are cases in which underprovision for one group allows for more profit by focusing on other groups. In the 'cream and park' which bedevils outsourced government services, contractors may cream off the easy and profitable cases whilst parking the more difficult, resource-intensive cases. So the incentive I'm considering is not the only one.

But I suspect it is the main one. Fundamentally, and imperfectly, marketised medicine responds to patient preferences. And, I will now argue, patients often prefer overprovision.

Sometimes this can be explicit, as in the case of buying a first-class ticket. I have in the USA been offered several tedious but painless investigations—including multiple expensive MRI scans—on the explicit basis that though my symptoms did not really justify the test, it would bring *some* benefit. Having paid my maximum insurance contribution for the year, there was no additional cost and thus 'no reason not to'. I accepted, of course. I have always loved a bargain. Overprovision particularly invites this kind of patient complicity *when the patient is not paying*: had I not made such use of my insurance policy that year, an MRI would have cost several hundred dollars. I had only have paid that much for the investigation given firm medical advice that it was worth the cost.

In England, payment does not typically enter the discussion. But patient satisfaction metrics are important. (If you have been to an NHS facility recently, you have most likely seen the 'Friends and Family Test' surveys.) Such metrics are one way patient opinions and reactions are taken into account, and patients like 'thorough' doctors. I find this claim generally plausible, and though it is hard to vindicate in general, more specific

instances are supported by evidence. For example, Ashworth et al. (2016) show that (in one study) higher levels of antibiotic prescriptions are correlated with patient satisfaction at the practice level.

In a context where supply underprovision is the norm, a preference for thoroughness can be rational for the patient. A doctor spending time and resources on a patient helps that patient get value for her money (and other 'outlays'). It is also psychologically reassuring—at the very least, the patient need not feel like a burdensome fool for coming in with something *obviously* not serious, if the doctor judges it worth a closer look.

The information asymmetry inherent to medical practice normally makes underprovision easier to detect than overprovision. Underprovision may engender untreated pain and illness, which are detectable by the patient and may lead to negative Friends and Family reviews, requests for a second opinion, damning coroner's reports, or simply patients moving to another doctor.

Conversely, overprovision may require medical expertise to detect, unless it is grossly excessive, such as prescribing codeine for a mild headache. Overprovision may be statistically detectable, or detectable by other doctors. But this is little help in a market, because it is normally harder for *the patient* to detect, and less likely to be resented or disliked by the patient. Even if the BMJ realises that Dr. Brown Bear's home visit was excessive, his patient may not, and may instead praise him for being caring and thorough like the family doctors of yore.

4 | Weighing Reasons in Clinical Judgement

I have argued that medical overprovision is harmful and thus morally bad (despite being less dramatic than underprovision). It harms both the patient involved and in the case of distributive overprovision, other patients who are unfairly deprived of resources.

I have also argued that it is likely to occur in many cases—the incentives are there. But given its badness, we have reason to minimise it. So what is the connection to payment structures?

The two classic tasks of the doctor are to diagnose and treat disease: to form beliefs and make decisions, often in the face of conflicting reasons or evidence. We do that all the time. The sky looks cloudy but the weather report says it will not rain. So I must form a belief: which is more trustworthy, the weather report or my eyes? And then I must make a practical decision: given my assessment of the weather, should I take the umbrella?

A doctor must first complete the epistemic task of forming a judgement about the patient's condition and clinical need. We generally think this judgement should be totally transparent to the patient, at least most of the time and with due care for the patient's emotional state. Then she must recommend one or other investigation or treatment or course of action (which might be nothing) to the patient. If the patient is unconscious or otherwise thought unable to make a decision, then she may have to make that decision on the patient's behalf.

Incentives *can* distort beliefs, but I will focus on the doctor's second, practical task, where they can really bite. She must weigh three sorts of practical reason. First, she has clinical reasons, determined by the outcome of her first (epistemic) task. That a patient's clinical need involves some kind of investigation or treatment is a clinical reason to recommend that course of action. That it would be painful with no expected benefit is a clinical reason not to.

The moral ideal is that only clinical reasons determine the doctor's recommendation. It is hard to see how doctors always recommending courses they sincerely judge to fit the patient's clinical need could be anything other than ideal, at least holding fixed the doctor's knowledge and ability, and so on.

Perhaps we might think that she should exercise some discretion when there's an effective cancer treatment costing £100,000 per QALY—and not funded by the NHS—and she knows her patient can not afford it. But those cases are quite marginal, and fall somewhere short of a moral *ideal*: it's not ideal if lifesaving treatment is unavailable for financial reasons. In any case, the psychological distress of knowing that there's an unaffordable treatment could be forced into the clinical reasons mould, albeit a little paternalistically, as the knowledge not being in the patient's best interests. Similarly, the first 'pro' of implicit—as opposed to explicit—rationing listed by (Gonzato 2022, 360) is that 'keeping patients in ignorance may be useful, because those who are unaware that a treatment exists but is not provided to them are less distressed'.

The doctor will almost always have *personal reasons* to act one way rather than another. The most obvious example is the financial, as in Dr. Brown Bear's case. But this is not intended to be a controversial point: some courses of action are simply more onerous or tedious. Those are *some* reasons not to take them—but that is not to say they should figure in her clinical deliberation.

It is morally better for doctors *not* to be swayed by personal reasons in clinical situations. Their being so swayed is morally bad. I can not think of an exception to this claim. Insofar as personal reasons lead the doctor to recommend or undertake a different course than she thinks required by her clinical reasons, this results in either overprovision, underprovision, or misprovision. As I have argued, these are ways in which medical treatment can go wrong.

Finally, doctors have *organisational reasons*. A course of action may have (e.g., financial or reputational) implications for the doctor's organisation. Some investigations and treatments will be more or less profitable: it is almost impossible that a business could precisely break even on every item individually, so even a hospital with an overall balanced budget will lose money on some things and make it on others. And a doctor may well know that if she recommends an expensive treatment regimen for a patient, it is more likely that the contract for treating such patients will be awarded to another ('more efficient') company in the next cycle.

Organisational reasons are not per se disreputable. A clinic that has gone bankrupt treating one patient will not be able to pay for

anyone else's treatment. It is quite reasonable for somebody to take them into account, just as rationing decisions must happen at some level. It is perfectly legitimate and necessary that organisational considerations are considered when determining what medical services are offered.

My focus is what happens in the consulting room, when a doctor is making decisions about a particular patient. There, I think *only* clinical reasons should figure in her deliberation.

5 | Against Markets

I thus think that a healthcare provider structure is morally criticisable insofar as it leaves doctors subject to non-clinical reasons, because it imposes a risk of them being so swayed. That risk is never totally eliminable. Even if Dr. Brown Bear was on a salary, he would have a reason (on one hand) to stay in his warm bed, and (on the other hand) to beef up his patient-satisfaction metrics. But corruption and temptation come in degrees.

The best way to avoid taking disreputable reasons into account is not to have them. As far as possible, doctors should have no personal reasons when making clinical decisions. Especially not financial reasons to, for example, recommend a particular investigation or treatment. A little more weakly, they should also have no organisational reasons.

I have argued that it is a problem when doctors are exposed to personal reasons. This grounds an argument for large provider organisations and against marketisation. In a large organisation without much exposure to market forces, the personal and organisational reasons facing doctors can be minimised. A salaried surgeon working for a large public hospital need not have any financial incentive to treat patients in some particular way. At the other end of the scale, his colleague who does not take a salary but bills the state or other insurer for each procedure has a direct financial incentive to maximise her income. A salary offers far less personal financial incentive than a fee for service does.

What does this have to do with markets? It would be true independent of the size of the organisation. Nevertheless, size matters. In smaller healthcare providers, organisational reasons translate to personal reasons. If the salaried doctor works not for a large hospital but is one of a consortium of three surgeons, then individual cases may determine the profitability of the consortium and thus his income this year. In general, if the consortium profits from certain courses of action then he has more or less direct financial reasons to keep the clinic profitable, or at the very least solvent.

The larger the organisation, the less individual clinical decisions affect its overall financial state. In the consulting room, this is a good thing. The NHS or even a single large hospital may marginally benefit financially from prescribing a slightly cheaper generic medication, but that is unlikely to give the salaried doctor a direct personal incentive simply because it is a drop in the ocean: the price difference is not likely to make a large hospital go bust.

But in a smaller and more fragmented healthcare economy—at the limit, an individual practitioner for whom organisational reasons *are* personal reasons—there may be more profit to be made from the branded drugs. Not for nothing are medical incentives to over-prescribe a central plot point in George Eliot's *Middlemarch*.

Even in a large NHS hospital, incentives can be constructed to give doctors personal reasons to act one way rather than another. A doctor's promotion prospects or annual bonus could be partially determined by the total costs of the treatments he recommends over the year. That would be a policy decision, an obviously bad one. But in a small enough provider—at the limit, a single doctor—the translation of organisational reasons into personal reasons is unavoidable. It would not be a policy decision but inherent to the structure of the business.

As I mentioned, organisational reasons are not always disreputable. Hospitals need to stay in business to ensure a fair distribution of healthcare if nothing else. I have argued that organisational reasons should not lead to personal reasons, however.

Though it perhaps can not be totally avoided, provider structures are unethical insofar as they force doctors to engage in a certain kind of *weighing deliberation* when recommending and offering care. We want doctors not to have to weigh organisational or personal reasons at all, but instead to act purely on the clinical ones.

Thus I suggest that the most effective solution is an *adversarial system*: in the clinical setting, the doctor has only clinical reasons to recommend one course of action over another. But these clinical judgements are constrained by an adversary, such as an insurance company or NHS commissioner, which determines whether treatments will be paid for. Now we come back to the parallel with bedside rationing. To avoid bedside profiteering, I have proposed a system akin to the 'gatekeeping' solution against bedside rationing where rationing is implemented 'in a fair process conducted at a higher level within the healthcare system, not by individual physicians', through the use of general rules (Lauridsen 2009, 318). If I have made plausible that such a gatekeeping or adversarial solution is morally optimal in the anti-profiteering case too, this bolsters its claim to be a sensible solution to the problems of rationing.

In such a system, the organisational considerations are not weighing reasons to be taken into account like any other, but more like fences that the doctor can try to get past. Thus, in dealing with individual patients, she always (insofar as it is possible) has only clinical reasons. The parenthetical caveat is not trivial: some patients and procedures are more annoying or time-consuming, some may be more professionally advantageous—might you get a publication out of this?—but in general, we especially wish to avoid direct *financial* incentives in the consulting room.

6 | Two Alternatives

I will now consider two alternative responses to the problem I have focused on. If the problem is that doctors have perverse

incentives, why not rely on professional codes and standards to ensure those incentives are not distorting, or why not work to align them with clinical need?

First, if a doctor is sufficiently indoctrinated into professional codes or virtues, she may be able to ignore any non-clinical reasons. And it might be argued that if doctors are not *in fact* swayed by non-clinical incentives, then there is no moral problem. I think this argument is overstated for two reasons. It is bad for the doctor to endure temptation, and exposure to such will damage the morally valuable doctor-patient relationship.

Exposing doctors to temptation exposes them to *moral risk*. This is bad for them on a purely hedonic level: resisting temptation is no fun. We should not focus purely on an (unscrupulous) caricature doctor out to fill his coffers and play golf. Many doctors are rich, but many are not. What if the cost of some investigation only slightly exceeds the expected benefit to the patient, and the fee will save your child from eviction? This is a much more troubling decision to face, and it would be better not to face it.

And how can one be sure one has not succumbed? Over a long career that is trying to resist professional and financial incentives to act in one way or another for several decades. Are you certain that you were *never* in fact swayed by them? Such certainty seems unrealistic, especially for a doctor dealing with difficult cases where the clinical reasons are murky and unclear.

The argument can be made at a statistical level. There are tens of thousands of doctors in the NHS, so if they routinely face personal reasons to act in one way rather than another, it is overwhelmingly likely that sometimes, some will act on those reasons, against the best interests of their patients.

Even if this *never* happens, we non-doctors can not be sure of that. Incentives in the presence of information asymmetry make it almost impossible to avoid the *suspicion* of medical corruption. And it would be bad for there to be such widespread suspicion. Doctors have multiple roles which make them particularly exposed to temptation. They diagnose patients, determine which investigations and treatments are called for, and potentially carry out those investigations or treatments, as well as communicating all of the relevant information and recommendations to the patient. All this may make it hard to dispel the suspicion of overprovision in any particular case. Doctors are the only people most of us trust to decide whether we need to be cut open, and then cut us open—and we pay them to do it.

Exposing professionals to moral risk and trusting them not to succumb to it is not a long-term solution—even if that trust is vindicated. I have argued that many NHS doctors are likely to be trusted more than other groups because they lack major personal incentives to act in one way or another. Trust between patient and doctor is morally valuable and should not be put at risk.

To return to the example of anonymous marking or journal reviews: we might well hope that as academics, our professional ethics prevent us from allowing non-academic facts about the student or author to influence our judgement. But an advantage of anonymity is that it spares us temptation, so we would not think that a turn to ‘open’ marking but simply trusting us not

to be biased is a sensible course, without very good reasons. Professional ethics and codes have an important role in medical and academic life, but we should not *rely* on them if we can instead bolster them.

Like the poor, incentives will always be with us. So perhaps instead of allowing their moral upshots to determine the structure of healthcare provision, perhaps a less drastic answer is to *align* the incentives with the patient's need? This is what happens when financial advisers are paid a percentage of the returns they earn for their clients. The thought then is that even if doctors act on their personal reasons, no harm is done because they lead to the same outcomes.

The British government has dipped its toes in these waters. For example, it claims in a discussion of capitation payments *for providers* that ‘allowing providers to share this money [ie, any capitation payments not spent on treatment] gives them an added incentive to keep patients in their target population healthy’.⁷

But attempting to always align incentives is simply very difficult. I have already mentioned ‘cream and park’ where private providers cream off the easy profitable cases and park the more difficult or costly ones. When we consider the complexity of healthcare—where outcomes are not as easily measured as a percentage return on a pot of money—it is almost impossible to imagine that the incentives align perfectly with patient interest, beyond either gross measures (minimising excess deaths) or tinkering around the edges (maximising Advice and Guidance). And of course overprovision is one way to minimise excess deaths...

There is a deeper moral question about whether we want aligned incentives in healthcare. I do not want my doctor determining my treatment even in part by whether it is the most profitable, even if the most profitable treatment is also the clinically appropriate one. I am in no position to assess whether that is in fact so, and the counterfactual question—what if some other treatment were more profitable?—will be quite unsettling, and so issues of trust once more arise.

Finally, we should not exaggerate the role aligned incentives play in financial services. The kind of adversarial approach that I have endorsed for medicine is also common where there really are conflicting reasons. In banking, for example, regulators often insist on structural protections such as so-called ‘Chinese walls’ or ‘ethical walls’ between divisions, to minimise the temptation and opportunity to do wrong.

7 | Conclusion

NHS doctors are lucky to be widely trusted. I have argued that this is at least partly because of the health service's structure. One of the NHS's main advantages is that you can (generally) trust that your doctor does not have financial incentives one way or another when it comes to your investigations and treatment. Since this is widely known, it grounds justified trust in doctors. Such trust is an ethical good, and something which should be taken into account when designing healthcare organisations, especially if we flirt with further marketisation.

I have argued that an adversarial system in which doctors do not face non-clinical reasons is morally optimal, at least with respect to the current issue. Marketisation introduces distinctive dangers.

I have not argued against NHS privatisation—the privatisation of NHS *provision*—but against its fragmentation and marketisation. Private organisations can be large and are (almost?) as capable of erecting an adversarial system as public ones are. But my argument does have implications for two common defences of privatisation.

First, one sometimes hears that it does not matter whether a doctor works for the state or for a private company, so long as the care is free. I have argued that this is dangerously simplistic: what kind of incentives she has may determine *what* care is provided, and even if not, the mere presence of such incentives imposes a moral risk on the doctor and may be corrosive to patient trust.

Second, a companions-in-guilt defence of privatisation is that the NHS buys in many products and services from private companies... should the production of chairs for hospitals be socialised, too? The argument is that it is irrational to draw some kind of moral line against doctors working for private companies, when those doctors are sitting in chairs and using equipment that is made by private companies. This defence fails, however, for two reasons. First, my argument appeals in part to distinctive features of medical care that are generally not shared with (for example) chair manufacturers. Second, insofar as those manufacturers were incentivised to overcharge or provide shoddy products to the NHS, this *would* be a—heavily defeasible!—argument for moving production in-house.

My argument has been *pro tanto*. Marketisation may have its merits, and they might even outweigh the ethical arguments I have made here. But the downside would be real—a risk of throwing away the precious good of doctors acting in the interests of patients and being *trusted* to do so.

Acknowledgements

For discussion and comments, I am indebted to Emma Borg, Robert Miller, Chris Newdick, Adam Pearce, an audience at the University of Reading, and two anonymous reviewers for Ratio.

Conflicts of Interest

The author declares no conflicts of interest.

Data Availability Statement

This paper uses purely theoretical arguments and, as such, there is no data.

Endnotes

¹ Bell (2017, 1).

² For a description of how the funding system works, see <https://www.england.nhs.uk/long-read/gp-contract/> from the NHS and the linked <https://www.kingsfund.org.uk/insight-and-analysis/long-reads/>

[gp-funding-and-contracts-explained](#) from the King's Fund health think tank.

³ See, for example, <https://www.pulsetoday.co.uk/news/practice-personal-finance/gps-to-be-paid-20-per-ag-request-in-bid-to-cut-hospital-waiting-lists/>.

⁴ <https://www.gov.uk/government/publications/the-nhs-constitution-for-england/the-nhs-constitution-for-england>.

⁵ See for example <https://www.nice.org.uk/process/pmg6/chapter/assessing-cost-effectiveness#economic-evidence-and-guideline-recommendations>.

⁶ Williams (1973, 240).

⁷ <https://www.gov.uk/guidance/capitation>, accessed September 2024.

References

- Ashworth, M., P. White, H. Jongsma, P. Schofield, and D. Armstrong. 2016. "Antibiotic Prescribing and Patient Satisfaction in Primary Care in England: Cross-Sectional Analysis of National Patient Survey Data and Prescribing Data." *British Journal of General Practice* 66, no. 642: e40–e46. <https://doi.org/10.3399/bjgp15X688105>.
- Bell, C. 2017. "Does Peppa Pig Encourage Inappropriate Use of Primary Care Resources?" *BMJ* 3597, no. December: j5397. <https://doi.org/10.1136/bmj.j5397>.
- Frith, L. 2018. "Organisational Ethics: A Solution to the Challenges of Markets in Healthcare?" In *Marketisation, Ethics and Healthcare*, edited by T. Feiler, J. Hordern, and A. Papanikitas. Routledge. <https://www.taylorfrancis.com/books/9781315186351>.
- Gonzato, O. 2022. "Bedside Rationing in Cancer Care: Patient Advocate Perspective." *Clinical Ethics* 17, no. 4: 358–362. <https://doi.org/10.1177/14777509211070486>.
- Lauridsen, S. 2009. "Administrative Gatekeeping—A Third Way Between Unrestricted Patient Advocacy and Bedside Rationing." *Bioethics* 23, no. 5: 311–320. <https://doi.org/10.1111/j.1467-8519.2008.00652.x>.
- Williams, B. 1973. "The Idea of Equality." In *Problems of the Self*, 230–249. Cambridge University Press. <https://doi.org/10.1017/CBO9780511621253.016>.